

Streaming approaches to error detection and trimming in the analysis of short sequencing reads

TBD

November 23, 2014

1 Introduction

K-mer spectral analysis is a powerful approach to error detection and correction in shotgun sequencing data that uses k-mer abundances to determine likely errors (cite Pevzner, Quake, khmer-counting paper). Approaches derived from spectral analysis can be very accurate: Zhang et al. (2014) show that spectral analysis is considerably more effective at finding errors than quality-based approaches (cite). However, spectral analysis is also very compute intensive; most implementations must count k-mers across entire sequencing data sets, which can be memory- or I/O-intensive for large data sets.

Streaming algorithms can offer improved algorithmic and computational efficiency in the analysis of large data sets (cite). Streaming algorithms typically examine the data only once, and scale in memory usage sublinearly with respect to the input data. Streaming algorithms have not been applied to k-mer spectral analysis of reads, although XXX (melsted).

Brown et al. (2012) introduced a streaming algorithm for normalizing k-mer abundance spectra, termed “digital normalization” (abbreviated as “diginorm”). This procedure estimates the k-mer coverage of each read in an online algorithm, by calculating the median k-mer abundance of the read given all previous reads; reads above a certain estimated coverage are set aside and their k-mers are not tracked. This algorithm is both online and *streaming* because it only collects k-mers in reads with a low estimated coverage; for the same reason, it is sublinear in memory for high coverage data sets. The net effect of diginorm is to reduce the data set size that must be considered for downstream processing, such as *de novo* assembly (cite trinity, elijah, etc.)

Here we develop a streaming algorithm for k-mer spectral analysis, based on digital normalization, that can detect and remove errors in sequencing reads. This algorithm operates in sublinear memory, and examines the data at most twice. The approach offers a general framework for making use of locus-specific graph saturation and could potentially be used for error correction, variant calling, and a streaming implementation of assembly. Moreover, it can be applied to data sets with variable coverage such as transcriptomes, metagenomes, and amplified genomic DNA.

2 Results

2.1 Coverage-normalized data can be used to locate errors in high-coverage shotgun sequencing data

To determine if digital normalization could be applied prior to k-mer spectral error detection, we first tried it on a simulated data set. We generated a synthetic data set from a simulated low-complexity genome (“simple genome”; see Methods for generation and Table XXX for data set details). We then applied digital normalization to a median 20-mer coverage of 20 ($k=20$, $C=20$).

The k-mer spectrum before and after digital normalization is shown in Figure 1. This also illustrates the concept underlying k-mer spectral error detection: in high coverage data sets, the *correct* k-mers are cleanly separated from the k-mers resulting from errors, and so a simple abundance cutoff can classify k-mers as correct or incorrect.




Figure 1: **K-mer spectrum of a simple artificial data set, before and after digital normalization.** The peaks at 1 represents erroneous k-mers resulting from (simulated) error; the peaks centered at 80 (raw) and 20 (diginorm) represent k-mers truly present in the genome which are present in many reads.

We next used k-mer counts from the downsampled read set to detect errors in the original read set. The algorithm is straightforward: we look for bases at the beginning or ends of low-abundance regions in each read. We used a “trusted k-mer” cutoff of $C_0 = 3$ as our abundance cutoff, below which we assumed k-mers were erroneous (see Methods). Of the 531 simulated reads from the simple genome containing one or more errors, predicted errors matched the known truth exactly for 443 of them (true positives), and 466 reads were correctly predicted

to contain no errors (true negatives). 0 reads were falsely predicted to have no errors (false negatives). The errors in 88 reads were miscalled – while the reads each had one or more errors, the positions were not correctly called – and three reads were incorrectly predicted to contain errors, leading to a total of 91 false positives. Using the above definitions, we calculated the prediction sensitivity to be 100% and the prediction specificity to be 83.0%.

This is similar to the results of spectral error detection using the unnormalized counts, which yielded 440 TP, 455 TN, 105 FP, and 0 FN, for a sensitivity of 100% and a specificity of 80.7%; the only difference in parameters for this analysis is that we used a cutoff of 10 rather than 3, to account for the shifted k-mer spectrum from diginorm.

@list in a table: matched exactly, etc. rundown. Columnar etc.

We next applied digital normalization and k-mer spectral error detection to an Illumina data set from *E. coli* MG1655 (cite Chitsaz). We mapped 5m untrimmed reads to the known *E. coli* MG1655 genome with bowtie1 (cite) and calculated mismatches between the reads and the genome as errors in the reads. This yielded 8.0m errors in 2.2m reads, for an overall error rate of 1.60%. Using these errors as ground truth, we found approximately 940,000 true positives, 2.8m true negatives, 895,000 false positives, and 356,000 false negatives. This yielded a sensitivity to errors of 72.4% and a specificity of 51.1%.

@add in unnormalized

2.2 Coverage-normalized data can be used to locate errors in variable coverage shotgun sequencing data

One of the drawbacks of spectral abundance analysis is that it cannot be applied to metagenomic or transcriptomic data sets, which frequently contain reads from both high-abundance and low-abundance molecules. This variability in coverage confounds naive spectral analysis for two reasons: first, errors in very high abundance regions can accumulate and increase over the threshold for trusted k-mers, thus appearing to be correct; and second, correct reads from low coverage regions yield k-mers below the trusted k-mer threshold and appear to be incorrect. In practice, therefore, error correction for metagenomic and transcriptome data must use other approaches.

If we apply diginorm to variable coverage data, we address the problem of accumulated errors from high-abundance regions by eliminating the majority of the reads from these regions; this is a form of error correction, detailed in the original diginorm paper (cite diginorm). We can also address the issue of low-abundance regions by ignoring reads that are at yet low coverage; here we can again use the median k-mer abundance to estimate which reads are too low-abundance to analyze (Figure XXX).

@graph with read abundance spectrum, showing which reads we will call errors in.

We generated two more synthetic data sets, “simple metagenome” and “simple mRNAseq,” which contain both high- and low-abundance species (see Table XXX for data set details). After generating synthetic reads with a 1% error

rate and applying digital normalization to $k=20/C=20$, we again applied spectral error detection using the normalized counts, but with the modification that only reads with a median k-mer abundance of 20 or greater were examined. For the simple metagenome data set, 2254 of 2347 reads (96.0%) met the coverage criterion. Of the 2347 reads total, the errors in 978 erroneous reads were called perfectly (TP) and 1125 of the reads with no errors were correctly called as error-free (TN). 68 reads were incorrectly determined to be error free (FN; including the reads that were too low coverage to be considered). Of the remaining 176 reads, 170 were miscalled (errors existed but were not exactly called) and 6 were incorrectly called as erroneous when they were in fact correct. We calculated the prediction sensitivity to be 93.5% and the prediction specificity to be 84.7%. For the simple mRNAseq data set, 524 of 568 reads (92.3%) met the coverage criterion, with 228 true positives, 250 true negatives, 29 false negatives (including the low coverage reads not examined), and 61 false positives, for a prediction sensitivity of 88.7% and prediction specificity of 78.9%.

@ real data / measurements - sherine, QP.

2.3 A streaming algorithm for error detection

Even with digital normalization, the spectral error detection approach outlined above is a 2-pass offline algorithm for any given data set - the first pass normalizes the read set and records the k-mer abundances, while the second pass examines the reads. A streaming approach could avoid some or all of a second pass, leading to greater efficiency.

Here we can make use of the redundancy of shotgun sequencing to avoid examining much of the data twice. Shotgun sequencing oversamples most regions - for example, for a 100x coverage genomic data set, we would expect 50% or more of the genome to be represented in more than 100 reads. This is a consequence of the Poisson-random sampling that underlies shotgun sequencing (cite... Waterman?) This oversampling provides an opportunity, however; if we regard the read data set as a stream of incoming data, some portions of the reference will more highly sampled earlier in the stream than others. For example, in mRNAseq, highly expressed transcripts will on average be highly sampled much earlier than low-expressed transcripts.

If highly sampled regions could be detected during iteration over the read data set, we could apply the same approaches used above to do error detection in a streaming fashion. Digital normalization provides such a method: by measuring coverage for each read against an online De Bruijn graph, reads that reach a specified coverage threshold can be examined for errors immediately. Those reads that had not yet reached high coverage could be set aside and re-examined later. This could result in significant runtime savings: for a genomic data set with 100x coverage, no more than 20% of the reads would need to be examined twice.

The conceptual idea is presented in Figure 2.

In Figure 3, we show diginorm-generated coverage saturation curves for both real and error-free simulated reads from *E. coli* MG1655. In both cases, we see




Figure 2: **Diagram of streaming error detection.** In a first pass over the read data, reads are loaded in until the graph locus to which they belong is saturated. From that point on, reads are examined for errors and not loaded into the graph. In a second pass, only the subset of reads loaded into the graph are examined for errors.

that by about 1m reads, the majority of reads comes from loci with an estimated sequencing depth of 20 or higher, and hence can be used for error analysis on the first pass through the data. Only those reads collected on the first pass need to be examined again.

When we apply this streaming approach to the simulated reads from the “simple genome”, we find identical numbers to the full two-pass approach: 443 TP, 466 TN, 91 FP, and 0 FN, for a sensitivity of 100% and a specificity of 83.0%. Likewise, for the “simple metagenome” and “simple mRNAseq” data sets, we obtain nearly identical and identical results, respectively, with the streaming approach; due to differences in the order in which reads are examined, the simple metagenome fails to detect one true positive and erroneously finds errors in three extra reads.

- @ do we put these numbers into a table? They’re basically identical.
- @ discuss streaming error detection results.
- @ apply to E. coli?




Figure 3: **Saturation curve of a real and a simulated *E. coli* read data set.** Reads are collected when they have an estimated coverage of less than 20; in the early phase ($< 1\text{m}$ reads), almost all reads are collected, but soon the majority of reads come from loci with an estimated sequencing depth of > 20 and are rejected.

2.4 A streaming algorithm for error trimming

Once errors can be *detected* with a streaming algorithm, it is simple to remove errors from reads by trimming the read at the first error. (While it is possible to split reads around errors rather than truncating them, this introduces complications in downstream read processing.)

On the “simple genome” with counts from the digitally normalized reads, this trimming approach eliminates 149 reads entirely due to a starting low-abundance k-mer, and truncates another 392 reads. Of the 100,000 bp in the simulated reads, 31,910 (31.9%) were removed by the trimming process. In exchange, trimming eliminated *all* of the errors, bringing the overall error rate from 0.63% to 0.00%.

For the simple metagenome we used the variable abundance approach described above and only trimmed reads with estimated coverage of 20 or higher. Here, of 2347 reads containing 234,700 bp, 314 reads (13.4%) were removed and 851 reads (36.3%) were trimmed, discarding a total of 74,321 bases (31.7%). Of 1451 errors total, all but 61 were eliminated, bringing the overall per-base error rate from 0.62% to 0.04%. The simple mRNAseq data set showed similar improvement: 83 of 568 reads were removed, and 208 were trimmed, removing 19,507 of 56,800 bases (34.34%). The initial error rate was 0.65% and the final error rate was 0.07%.

Applying the streaming error trimming to the *E. coli* MG1655 data set used

in section XXX, we trimmed 2.0m reads and removed nearly 122,000 reads entirely. Of 7.6m errors, all but 198,000 were removed, bringing the error rate from 1.60% to 0.5%. Trimming discarded 64 Mbp of the original 500 Mbp (13.0%).

@ mrna @ real metagenome

2.5 Illumina error rates and error profiles can be determined from a small sample of sequencing data

With Illumina sequencing, average and per-position error rates may vary between sequencing runs, but are typically systematic within a run (cite?). Per-position error rates are caused by fluidics etc.

We can use the approaches described above to calculate systematic error profiles for shotgun sequencing data across entire data sets, but the variable abundance approach developed above can also be applied to *subsets* of Illumina data. The essential idea is to consume reads until sufficient data has been collected to calculate error rates, and then to calculate those error rates for the new reads based on the k-mer abundances from the old reads. This can also be done in one pass for data sets with sufficiently high coverage data: as shown above (Figure XXX), more than half of the reads will have sufficient coverage to call errors by the time 10% of the data set has been consumed.

We first simulated a set of reads from the simple genome with errors only at even positions (0, 2, 4, etc.), and called errors in these reads using the single-pass algorithm described above (see Methods for implementation details). We also separately called errors by mapping with bowtie and examining mismatches. The results are shown in Figure YYY; the difference in predicted vs actual mismatch profiles has an average of 0.04%, with a variance of 0.0002% across positions. The predicted per-base error rate is 0.57%, while the true error rate is 0.65%; this underestimate may be due to ignoring reads with multiple errors in them that present as too low-coverage to assess.

@@show figure with shaded color indicating difference

We next applied to ecoli data.@@

@@error rates? apply regression?

3 Discussion

3.1 Digital normalization enables k-mer spectral error detection in variable abundance shotgun data

Reference-free error detection and correction in shotgun data from metagenomic, transcriptomic, and amplified samples typically requires specialized approaches (cite, discuss). Here we demonstrate that after digital normalization, k-mer abundance spectral approaches can be applied to reads estimated to have high coverage. This should generalize to any error detection, trimming, and counting approaches that rely on k-mer abundance.

One significant disadvantage of the variable abundance approach is that low-coverage reads cannot be trimmed because we do not know whether they are from low-abundance molecules or are highly erroneous. This must be taken into account for downstream analysis; for example, while assemblers must already ignore or correct these erroneous reads, quantitation approaches using e.g. mapping should not be applied to the trimmed data.

(@This is where applying Quake would be a useful demonstration.)

3.2 Read coverage estimates enable streaming few-pass approaches to error detection and trimming

K-mer spectral error detection and trimming approaches require two complete passes through the data - one for calculating k-mer abundances, and one for trimming reads. Here we develop an approach to streaming error detection and trimming that relies on uneven coverage saturation as the read data set is traversed. Again, this approach should be general.

3.3 Empirical error profiles can be calculated on subsets of Illumina shotgun data

Useful for cores etc: quick evaluation of sequencing quality, regardless of origin.

3.4 Concluding thoughts

Primary practical - The kind of trimming we implement is only useful for low memory assembly of variable coverage data sets, but we have plenty.

Time- and memory-efficient error profile calculation is practically useful for large sequencing cores.

Dealing with variants is outside the scope of this work, but 50/50 variants should not be trimmed. What about repeats?

Streaming methods for error correction and variant calling, more generally.

Theory needed for thresholding, but, in theory, thresholds should be static. (We can demonstrate that there is little variation.?)